EXPERIMENT-6

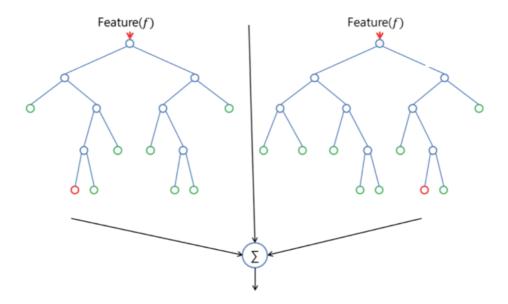
Problem Statement

Develop a machine learning method to predict stock price based on past price variation.

Algorithm

Random forest is a flexible, easy to use <u>machine learning algorithm</u> that produces, even without hyper-parameter tuning, a great result most of the time. It is also one of the most used algorithms, because of its simplicity and diversity (it can be used for both classification and regression tasks). In this post we'll learn how the random forest algorithm works, how it differs from other algorithms and how to use it.

One big advantage of random forest is that it can be used for both classification and regression problems, which form the majority of current machine learning systems. Let's look at random forest in classification, since classification is sometimes considered the building block of machine learning. Below you can see how a random forest would look like with two trees:



Random forest has nearly the same hyperparameters as a decision tree or a bagging classifier. Fortunately, there's no need to combine a decision tree with a bagging classifier because you can easily use the classifier-class of random forest. With random forest, you can also deal with regression tasks by using the algorithm's regressor.

Program Snippet

```
In [1]: import numpy as np
import pandas as pd

In [2]: df= pd.read_csv('stock.csv')
df
```

Out[2]:

	Date	Open	High	Low	Last	Close	Total Trade Quantity	Turnover (Lacs)
0	2018-10-08	208.00	222.25	206.85	216.00	215.15	4642146	10062.83
1	2018-10-05	217.00	218.60	205.90	210.25	209.20	3519515	7407.06
2	2018-10-04	223.50	227.80	216.15	217.25	218.20	1728786	3815.79
3	2018-10-03	230.00	237.50	225.75	226.45	227.60	1708590	3960.27
4	2018-10-01	234.55	234.60	221.05	230.30	230.90	1534749	3486.05
1230	2013-10-14	160.85	161.45	157.70	159.30	159.45	1281419	2039.09
1231	2013-10-11	161.15	163.45	159.00	159.80	160.05	1880046	3030.76
1232	2013-10-10	156.00	160.80	155.85	160.30	160.15	3124853	4978.80
1233	2013-10-09	155.70	158.20	154.15	155.30	155.55	2049580	3204.49
1234	2013-10-08	157.00	157.80	155.20	155.80	155.80	1720413	2688.94

1235 rows × 8 columns

In [3]: df.head()

Out[3]:

	Date	Open	High	Low	Last	Close	Total Trade Quantity	Turnover (Lacs)
0	2018-10-08	208.00	222.25	206.85	216.00	215.15	4642146	10062.83
1	2018-10-05	217.00	218.60	205.90	210.25	209.20	3519515	7407.06
2	2018-10-04	223.50	227.80	216.15	217.25	218.20	1728786	3815.79
3	2018-10-03	230.00	237.50	225.75	226.45	227.60	1708590	3960.27
4	2018-10-01	234.55	234.60	221.05	230.30	230.90	1534749	3486.05

In [4]: df.tail()

Out[4]:

	Date	Open	High	Low	Last	Close	Total Trade Quantity	Turnover (Lacs)
1230	2013-10-14	160.85	161.45	157.70	159.3	159.45	1281419	2039.09
1231	2013-10-11	161.15	163.45	159.00	159.8	160.05	1880046	3030.76
1232	2013-10-10	156.00	160.80	155.85	160.3	160.15	3124853	4978.80
1233	2013-10-09	155.70	158.20	154.15	155.3	155.55	2049580	3204.49
1234	2013-10-08	157.00	157.80	155.20	155.8	155.80	1720413	2688.94

```
In [5]: df.info()
        <class 'pandas.core.frame.DataFrame'>
        RangeIndex: 1235 entries, 0 to 1234
        Data columns (total 8 columns):
         # Column
                                   Non-Null Count Dtype
                                   -----
             -----
                                   1235 non-null object
         0
            Date
                                  1235 non-null float64
         1
             Open
         2
             High
         3
            Low
         4
            Last
         5
            Close
            Total Trade Quantity 1235 non-null int64
Turnover (Lacs) 1235 non-null float64
         6
         7 Turnover (Lacs)
        dtypes: float64(6), int64(1), object(1)
        memory usage: 77.3+ KB
In [6]: df.shape
Out[6]: (1235, 8)
In [7]: df.columns.values
```

In [8]: df.corr()

Out[8]:

Open	High	Low	Last	Close	Total Trade Quantity	Turnover (Lacs)
1.000000	0.998956	0.998776	0.997662	0.997704	0.367503	0.587026
0.998956	1.000000	0.998728	0.999130	0.999159	0.388798	0.605907
0.998776	0.998728	1.000000	0.999008	0.999065	0.361695	0.582446
0.997662	0.999130	0.999008	1.000000	0.999963	0.381269	0.599575
0.997704	0.999159	0.999065	0.999963	1.000000	0.380801	0.599155
0.367503	0.388798	0.361695	0.381269	0.380801	1.000000	0.941976
0.587026	0.605907	0.582446	0.599575	0.599155	0.941976	1.000000
	1.000000 0.998956 0.998776 0.997662 0.997704 0.367503	1.000000 0.998956 0.998956 1.000000 0.998776 0.998728 0.997662 0.999130 0.997704 0.999159	1.000000 0.998956 0.998776 0.998956 1.000000 0.998728 0.998776 0.998728 1.000000 0.997662 0.999130 0.999008 0.997704 0.999159 0.999065 0.367503 0.388798 0.361695	1.000000 0.998956 0.998776 0.997662 0.998956 1.000000 0.998728 0.999130 0.998776 0.998728 1.000000 0.999008 0.997662 0.999130 0.999008 1.000000 0.997704 0.999159 0.999065 0.999963 0.367503 0.388798 0.361695 0.381269	1.000000 0.998956 0.998776 0.997662 0.997704 0.998956 1.000000 0.998728 0.999130 0.999159 0.998776 0.998728 1.000000 0.999008 0.999065 0.997662 0.999130 0.999008 1.000000 0.999963 0.997704 0.999159 0.999065 0.999963 1.000000 0.367503 0.388798 0.361695 0.381269 0.380801	1.000000 0.998956 0.998776 0.997662 0.997704 0.367503 0.998956 1.000000 0.998728 0.999130 0.999159 0.388798 0.998776 0.998728 1.000000 0.999008 0.999065 0.361695 0.997662 0.999130 0.999008 1.000000 0.999963 0.381269 0.997704 0.999159 0.999065 0.999963 1.000000 0.380801 0.367503 0.388798 0.361695 0.381269 0.380801 1.000000

```
In [9]: df['Date']= pd.to_datetime(df.Date, format= '%Y-%m-%d')
    df.index = df['Date']
    df
```

Out[9]:

	Date	Open	High	Low	Last	Close	Total Trade Quantity	Turnover (Lacs)
Date								
2018-10-08	2018-10-08	208.00	222.25	206.85	216.00	215.15	4642146	10062.83
2018-10-05	2018-10-05	217.00	218.60	205.90	210.25	209.20	3519515	7407.06
2018-10-04	2018-10-04	223.50	227.80	216.15	217.25	218.20	1728786	3815.79
2018-10-03	2018-10-03	230.00	237.50	225.75	226.45	227.60	1708590	3960.27
2018-10-01	2018-10-01	234.55	234.60	221.05	230.30	230.90	1534749	3486.05
2013-10-14	2013-10-14	160.85	161.45	157.70	159.30	159.45	1281419	2039.09
2013-10-11	2013-10-11	161.15	163.45	159.00	159.80	160.05	1880046	3030.76
2013-10-10	2013-10-10	156.00	160.80	155.85	160.30	160.15	3124853	4978.80
2013-10-09	2013-10-09	155.70	158.20	154.15	155.30	155.55	2049580	3204.49
2013-10-08	2013-10-08	157.00	157.80	155.20	155.80	155.80	1720413	2688.94

1235 rows × 8 columns

```
In [10]: import matplotlib.pyplot as plt plt.figure(figsize=(16,8)) plt.plot(df['close'],label='closeHistory')

Out[10]: [<matplotlib.lines.Line2D at 0x1507bbbd0a0>]

250

250

150

100
```

Long Short Term Memory (LSTM)

LSTMs are widely used for sequence prediction problems and have proven to be extremely effective. The reason they work so well is because LSTM is able to store past information that is important, and forget the information that is not. LSTM has three gates:

The input gate: The input gate adds information to the cell state

The forget gate: It removes the information that is no longer required by the model

The output gate: Output Gate at LSTM selects the information to be shown as output

```
In [11]: data= df.sort_index(ascending=True, axis=0)
    new_data = pd.DataFrame(index=range(0,len(df)), columns=['Date','Close'])
    new_data
```

Out[11]:

	Date	Close
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN
1230	NaN	NaN
1231	NaN	NaN
1232	NaN	NaN
1233	NaN	NaN
1234	NaN	NaN

1235 rows × 2 columns

```
In [12]: for i in range(0,len(data)):
    new_data['Date'][i]= data['Date'][i]
    new_data['Close'][i]= data['Close'] [i]
    new_data
```

Out[12]:

	Date	Close
0	2013-10-08 00:00:00	155.8
1	2013-10-09 00:00:00	155.55
2	2013-10-10 00:00:00	160.15
3	2013-10-11 00:00:00	160.05
4	2013-10-14 00:00:00	159.45
1230	2018-10-01 00:00:00	230.9
1231	2018-10-03 00:00:00	227.6
1232	2018-10-04 00:00:00	218.2
1233	2018-10-05 00:00:00	209.2
1234	2018-10-08 00:00:00	215.15

1235 rows × 2 columns

```
In [13]: new_data.index = new_data.Date
         new_data.drop('Date', axis=1, inplace=True)
In [14]: dataset = new_data.values
         train = dataset[0:987,:]
         valid = dataset[987:,:]
In [15]: from sklearn.preprocessing import MinMaxScaler
         from keras.models import Sequential
         from keras.layers import Dense, Dropout, LSTM
In [16]: scaler = MinMaxScaler(feature_range=(0, 1))
         scaled_data = scaler.fit_transform(dataset)
In [17]: x_train, y_train = [], []
         for i in range(60,len(train)):
             x train.append(scaled data[i-60:i,0])
             y_train.append(scaled_data[i,0])
         x_train, y_train = np.array(x_train), np.array(y_train)
         x_train = np.reshape(x_train, (x_train.shape[0],x_train.shape[1],1))
```

```
In [18]: model = Sequential()
              model.add(LSTM(units=50, return sequences=True, input shape=(x train.shape[1],1)))
              model.add(LSTM(units=50))
              model.add(Dense(1))
              model.compile(loss='mean_squared_error', optimizer='adam')
              model.fit(x_train, y_train, epochs=1, batch_size=1, verbose=2)
              927/927 - 50s - loss: 0.0012
 Out[18]: <keras.callbacks.History at 0x1504549d1f0>
  In [19]: inputs = new_data[len(new_data) - len(valid) - 60:].values
              inputs = inputs.reshape(-1,1)
              inputs = scaler.transform(inputs)
  In [20]: X_test = []
              for i in range(60,inputs.shape[0]):
                  X_test.append(inputs[i-60:i,0])
              X test = np.array(X test)
              X_test = np.reshape(X_test, (X_test.shape[0],X_test.shape[1],1))
              closing_price = model.predict(X_test)
              closing_price = scaler.inverse_transform(closing_price)
In [21]: train = new_data[:987]
  valid = new_data[987:]
  valid['Predictions'] = closing_price
  plt.plot(train['close'])
  plt.plot(valid[['close', 'Predictions']])
         <ipython-input-21-f5cac910112e>:3: SettingWithCopyWarning:
         A value is trying to be set on a copy of a slice from a DataFrame. Try using .loc[row_indexer,col_indexer] = value instead
         See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-ve
          valid['Predictions'] = closing_price
Out[21]: [<matplotlib.lines.Line2D at 0x1504c112c40>, <matplotlib.lines.Line2D at 0x1504c112c70>]
          300
          250
          100
               2014
                       2015
                               2016
                                      2017
                                              2018
```

Github link-

https://github.com/AkSingh03/MACHINE LEARNING